

Signed Integers:

Signed Integers:

Faster *and* Correcter

Background

Background

- [unsigned: A Guideline for Better Code](#) (Jon Kalb, CppCon 2016)

Background

- [unsigned: A Guideline for Better Code](#) (Jon Kalb, CppCon 2016)
- [Garbage In, Garbage Out: Arguing about Undefined Behavior...](#) (Chandler Carruth, CppCon 2016)

Background

- [unsigned: A Guideline for Better Code](#) (Jon Kalb, CppCon 2016)
- [Garbage In, Garbage Out: Arguing about Undefined Behavior...](#) (Chandler Carruth, CppCon 2016)
- [Random thoughts on programming languages, compilers, operating systems, etc.](#) (Krister Walfridsson's blog, 2016-02-21)

Krister Walfridsson's blog

Random thoughts on programming languages, compilers, operating systems, etc.

Sunday, February 21, 2016

How undefined signed overflow enables optimizations in GCC

Signed integers are not allowed to overflow in C and C++, and this helps compilers generate better code. I was interested in how GCC is taking advantage of this, and here are my findings.¹

Signed integer expression simplification

The nice property of overflow being undefined is that signed integer operations works as in normal mathematics — you can cancel out values so that $(x*10)/5$ simplifies to $x*2$, or $(x+1) < (y+3)$ simplifies to $x < (y+2)$. Increasing a value always makes it larger, so $x < (x+1)$ is always true.

GCC iterates over the IR (the compiler's Internal Representation of the program), and does the following transformations (x , and y are signed integers, c , $c1$, and $c2$ are positive constants, and cmp is a comparison operator. I have only listed the transformations for positive constants, but GCC handles negative constants too in the obvious way)

- Eliminate multiplication in comparison with 0

$$(x * c) \text{ cmp } 0 \quad \rightarrow \quad x \text{ cmp } 0$$

- Eliminate division after multiplication

$$(x * c1) / c2 \quad \rightarrow \quad x * (c1 / c2) \text{ if } c1 \text{ is divisible by } c2$$

- Eliminate negation

$$(-x) / (-y) \quad \rightarrow \quad x / y$$

- Simplify comparisons that are always true or false

Contracts in C++

Contracts in C++

```
void do_something(Object const& thing);
```

Contracts in C++

```
void do_something(Object const& thing);
```

Question : Is const liberating or restrictive?

Contracts in C++

```
void do_something(Object const& thing);
```

Question : Is const liberating or restrictive?

Answer : It depends.

Contracts in C++

```
void do_something(Object const& thing);
```

Question : Is const liberating or restrictive?

Answer : It depends.

Implementor : "It's restrictive!"

Contracts in C++

```
void do_something(Object const& thing);
```

Question : Is `const` liberating or restrictive?

Answer : It depends.

Implementor : "It's restrictive!"

User : "It's liberating!"

Narrow Contracts in C++

```
void visit_something(std::function<void(Object const&)>);
```

Narrow Contracts in C++

```
void visit_something(std::function<void(Object const&)>);
```

```
visit_something([](Object const& thing) {  
    const_cast<Object&>(thing) = Object{}; // bad idea  
});
```

Narrow Contracts in C++

```
void visit_something(std::function<void(Object const&)>);
```

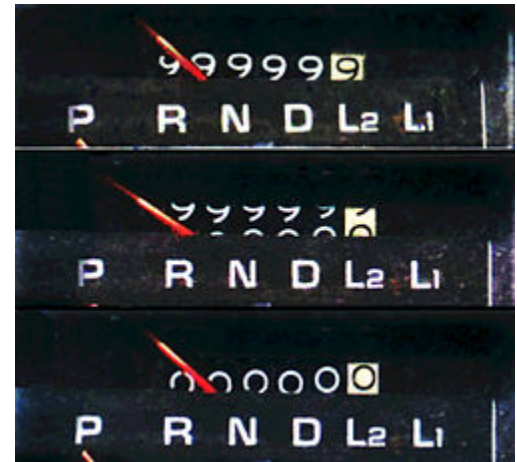
```
visit_something([](Object const& thing) {  
    const_cast<Object&>(thing) = Object{}; // bad idea  
});
```

How does \uparrow differ from \downarrow ?

```
auto i = numeric_limits<int>::max();  
++ i; // very bad idea
```


Defined but Wrong

```
// error: static_assert expression is not an integral constant expression  
static_assert(numeric_limits<signed>::max()+2);  
  
// totally fine  
static_assert(numeric_limits<unsigned>::max()+2);
```



If you're arguing...

C++ source #1 ×

```
1 bool example4a(T x) {  
2     T c = 1;  
3     return x + c < x;  
4 }  
5  
6 bool example4b(T x) {  
7     T c = 1;  
8     return x + c <= x;  
9 }
```

x86-64 gcc 7.1 (Editor #1, Compiler #1) ×

x86-64 gcc 7.1 ▾

-O3 -DT=signed

A ▾ 11010 .LX0: .text // \s+ Intel Demangle

```
1 example4a(int):  
2     xor     eax, eax  
3     ret  
4 example4b(int):  
5     xor     eax, eax  
6     ret
```

x86-64 gcc 7.1 (Editor #1, Compiler #2) ×

x86-64 gcc 7.1 ▾

-O3 -DT=unsigned

A ▾ 11010 .LX0: .text // \s+ Intel Demangle

```
1 example4a(unsigned int):  
2     cmp     edi, -1  
3     sete    al  
4     ret  
5 example4b(unsigned int):  
6     cmp     edi, -1  
7     sete    al  
8     ret
```

C++ source #1 ×

```
1 example2(T x) {  
2     return (x * 14) / 7;  
3 }
```

x86-64 gcc 7.1 (Editor #1, Compiler #1) ×

x86-64 gcc 7.1 ▾

-O3 -DT=signed

A ▾ 11010 .LX0: .text // \s+ Intel Demangle

```
1 example2(int):  
2     lea     eax, [rdi+rdi]  
3     ret
```

x86-64 gcc 7.1 (Editor #1, Compiler #2) ×

x86-64 gcc 7.1 ▾

-O3 -DT=unsigned

A ▾ 11010 .LX0: .text // \s+ Intel Demangle

```
1 example2(unsigned int):  
2     imul    edi, edi, 14  
3     mov     edx, 613566757  
4     mov     eax, edi  
5     mul     edx  
6     sub     edi, edx  
7     shr     edi  
8     lea     eax, [rdx+rdi]  
9     shr     eax, 2  
10    ret
```

C++ source #1 ×

A ▾   

```
1 bool example1(T x) {  
2     T c = 5;  
3     return (x * c) > 0;  
4 }
```


x86-64 gcc 7.1 (Editor #1, Compiler #1) ×

x86-64 gcc 7.1 ▾

-O3 -DT=signed

A ▾ 11010 .LX0: .text // \s+ Intel Demangle   




```
1 example1(int):  
2     test    edi, edi  
3     setg    al  
4     ret
```

 g++ (GCC-Explorer-Build) 7.1.0- cached

x86-64 gcc 7.1 (Editor #1, Compiler #2) ×

x86-64 gcc 7.1 ▾

-O3 -DT=unsigned

A ▾ 11010 .LX0: .text // \s+ Intel Demangle   

```
1 example1(unsigned int):  
2     lea     eax, [rdi+rdi*4]  
3     test    eax, eax  
4     setne   al  
5     ret
```

 g++ (GCC-Explorer-Build) 7.1.0- cached

C++ source #1 ×

```
1 #include <cstdlib>
2 T example9d(T x) {
3     if (x > 0) {
4         return abs(x);
5     }
6 }
```

x86-64 gcc 7.1 (Editor #1, Compiler #1) ×

x86-64 gcc 7.1 ▾

-O3 -DT=signed

```
A- 11010 .LX0: .text // \s+ Intel Demangle
1 example9d(int):
2     test    edi, edi
3     jle     .L2
4     mov     eax, edi
5     ret
6 .L2:
7     rep ret
```

g++ (GCC-Explorer-Build) 7.1.0- cached

x86-64 gcc 7.1 (Editor #1, Compiler #2) ×

x86-64 gcc 7.1 ▾

-O3 -DT=unsigned

```
A- 11010 .LX0: .text // \s+ Intel Demangle
1 example9d(unsigned int):
2     test    edi, edi
3     je      .L2
4     mov     eax, edi
5     sar     eax, 31
6     xor     edi, eax
7     sub     edi, eax
8     mov     eax, edi
9     ret
10 .L2:
11     rep ret
```

g++ (GCC-Explorer-Build) 7.1.0- cached

C++ source #1 ×



```
1 bool example8(float * a, T i) {  
2     float * e0 = a+i;  
3     float * e1 = a+(i+1);  
4     return e1 - e0 == 1;  
5 }
```

[Full-screen Snip](#)

x86-64 gcc 7.1 (Editor #1, Compiler #1) ×

x86-64 gcc 7.1 ▾

-O3 -DT=signed



11010

.LX0: .text //

\s+

Intel

Demangle



```
1 example8(float*, int):  
2     mov     eax, 1  
3     ret
```

g++ (GCC-Explorer-Build) 7.1.0- cached

x86-64 gcc 7.1 (Editor #1, Compiler #2) ×

x86-64 gcc 7.1 ▾

-O3 -DT=unsigned



11010

.LX0: .text //

\s+

Intel

Demangle



```
1 example8(float*, unsigned int):  
2     lea     eax, [rsi+1]  
3     mov     esi, esi  
4     sub     rax, rsi  
5     sal     rax, 2  
6     cmp     rax, 4  
7     sete    al  
8     ret
```

g++ (GCC-Explorer-Build) 7.1.0- cached

Compiler Explorer

C++EditorDiff ViewMore

ShareOther

C++ source #1

A-

```
1 T example10(T m) {
2     int sum = 0;
3     for (short T i = 0; i <= m; ++ i) {
4         sum += i;
5     }
6     return sum;
7 }
```

x86-64 gcc 7.1 (Editor #1, Compiler #1)

x86-64 gcc 7.1 (Editor #2, Compiler #2)

x86-64 gcc 7.1

-O3 -DT=signed

A-

11010

LX0

text

//

/s+

Intel

Demangle

```
1 example10(int):
2     test    edi, edi
3     js      .L4
4     xor     ecx, ecx
5     xor     eax, eax
6     xor     edx, edx
7 .L3:
8     add     ecx, 1
9     add     eax, edx
10    movzx   edx, cx
11    cmp     edx, edi
12    jle     .L3
13    rep     ret
14 .L4:
15    xor     eax, eax
16    ret
```

Full-screen Snip

x86-64 gcc 7.1

-O3 -DT=unsigned

A-

11010

LX0

text

//

/s+

Intel

Demangle

```
1 example10(unsigned int):
2     lea     eax, [rdi-12]
3     cmp     eax, 65522
4     ja      .L6
5     lea     ecx, [rdi+1]
6     pxor    xmm0, xmm0
7     pxor    xmm2, xmm2
8     xor     eax, eax
9     mov     ecx, ecx
10    mov     xmm1, XMMWORD PTR .LC0[rip]
11    shr     edx, 3
12    movdqa   xmm4, XMMWORD PTR .LC1[rip]
13 .L3:
14    movdqa   xmm3, xmm1
15    add     eax, 1
16    cmp     eax, edx
17    punpcklwd    xmm3, xmm2
18    padd     xmm0, xmm3
19    movdqa   xmm3, xmm1
20    paddw    xmm1, xmm4
21    punpckhwd    xmm3, xmm2
22    padd     xmm0, xmm3
23    jb      .L3
24    movdqa   xmm1, xmm0
25    mov     edx, ecx
26    and     edx, -8
27    psrldq   xmm1, 8
28    padd     xmm0, xmm1
29    movdqa   xmm1, xmm0
30    cmp     ecx, edx
31    psrldq   xmm1, 4
32    padd     xmm0, xmm1
33    movd     eax, xmm0
34    je      .L5
35    lea     ecx, [rdx+1]
36    add     eax, edx
37    movzx   ecx, cx
38    cmp     edi, ecx
39    jb      .L5
40    add     eax, ecx
41    lea     ecx, [rdx+2]
42    movzx   ecx, cx
43    cmp     edi, ecx
44    jb      .L5
45    add     eax, ecx
46    lea     ecx, [rdx+3]
47    movzx   ecx, cx
48    cmp     edi, ecx
49    jb      .L5
50    add     eax, ecx
51    lea     ecx, [rdx+4]
52    movzx   ecx, cx
53    cmp     edi, ecx
54    jb      .L5
55    add     eax, ecx
56    lea     ecx, [rdx+5]
57    movzx   ecx, cx
58    cmp     edi, ecx
59    jb      .L5
```

g++ (GCC-Explorer-Build) 7.1.0-cached

g++ (GCC-Explorer-Build) 7.1.0-cached

Notes

- My knowledge of assembler gets a little hazy after 6502.
- Fewer instructions does *not* mean faster code.
- Some 'fine tuning' of details to make some of the results differ between signed/unsigned.
- I'll bet some unsigned optimizations are pretty impressive too so...
- Go and find some good counter-examples!

Thank You!

Special Thanks to Krister Walfridsson & Bay Area ACCU

c/c++ == 1

(signed) John McFarlane

[@JSAMcFarlane](https://twitter.com/JSAMcFarlane)